

ASSIGNMENT 4: THEORY OF GENERATIVE MODELS [IFT6135]

JOSEPH D. VIVIANO

1. REPARAMETERIZATION TRICK OF VARIATIONAL AUTOENCODERS

1.1. Transformation of Gaussian Noise.

1.2. Encoders vs. Mean Fields.

2. IMPORTANCE WEIGHTED AUTOENCODER

2.1. IWLB as a Lower Bound on log Likelihood.

2.2. IWLB with k=2 is tighter than ELBO with k=1.

3. MAXIMUM LIKELIHOOD FOR GENERATIVE ADVERSARIAL NETWORKS

The original GAN objective can be written as:

$$\max_D \mathbf{E}_{p_D(x)} [\log D(x)] + \mathbf{E}_{p_G} [\log(1 - D(G(z)))]; \quad \max_G \mathbf{E}_{p_G} [\log D(G(z))]$$

Note that the definition of an optimal discriminator using this notation is:

$$(3.1) \quad D^*(\mathbf{x}) = \frac{p_D(\mathbf{x})}{p_G(\mathbf{x}) + p_D(\mathbf{x})}$$

The goal here is to find the maximum likelihood objective (cost function) instead of the negative log likelihood objective to apply to samples coming from the generator $\mathbf{E}_{p_G} [f(D(G(z)))]$, where $f(D(G(z)))$ must be found.

As a reminder, the maximum likelihood estimate in this case would be:

$$\hat{\theta} \in \{\arg \max_{\theta \in \Theta} \{(\theta; x)\}\}$$

Each step of learning in a GAN consists of reducing the expectation $f(x)$ run on a bunch of samples pulled from generator G , which we express as

$$(3.2) \quad \mathbf{E}_{x \sim p_G} f(x)$$

here, p_G represents a sample pulled from the probability distribution generated by G . First, let's take partial derivative with respect to the weights θ on a sample $x \sim p_G$ from the generator and represent it as an integral. Then we applied Leibniz's rule, and sub in the identity $\frac{\partial}{\partial \theta} p_G(x) = p_G(x) \frac{\partial}{\partial \theta} \log p_G(x)$:

$$(3.3) \quad \begin{aligned} \frac{\partial}{\partial \theta} \mathbf{E}_{x \sim p_G} f(x) &= \int f(x) \frac{\partial}{\partial \theta} p_G(x) \\ &= \int f(x) \frac{\partial}{\partial \theta} p_G(x) dx \\ &= \int f(x) p_G(x) \frac{\partial}{\partial \theta} \log p_G(x) \end{aligned}$$

While tells us that we can express the aforementioned expectation (3.2) as:

$$(3.4) \quad \mathbf{E}_{x \sim p_G} f(x) \frac{\partial}{\partial \theta} \log p_G(x)$$

This tells us the maximum likelihood can be found given:

$$(3.5) \quad f(x) = -\frac{p_D(x)}{p_G(x)}$$

Now if we assume that our optimal discriminator $D^*(\mathbf{x})$ from (3.1) is the logistic sigmoid $\sigma(a(x))$ then

$$(3.6) \quad \sigma(a(x)) = \frac{p_D(\mathbf{x})}{p_G(\mathbf{x}) + p_D(\mathbf{x})}$$

And it follows that

$$(3.7) \quad f(x) = -\exp(a(x))$$

which is our function f such that the objective corresponds to maximum likelihood.

UNIVERSITÉ DE MONTRÉAL

E-mail address: `joseph@viviano.ca`